


HD95: 2D vs 2.5D context, ENet and residual-encoder U-Net

4-fold cross-validation × 3 seeds, 40 patients. 40-patient SEGTHOR train set (corrected labels), 4 folds × 3 seeds, median spacing + HU window + z-score + ROI crop, CE + soft Dice loss. U-Net has deep supervision. 2.5D = the slice and its ±k neighbours as 2k+1 input channels, centre slice predicted



HD95 mm, 3D on the original CT grid, best checkpoint; each run's value is the mean over the patients of its validation fold. All organs = mean of the 4 organs per run. Y-axes differ per organ.