


HD95: 2.5D context (3 slices) vs the 2D baseline, residual-encoder U-Net

4-fold cross-validation × 3 seeds, 40 patients. Residual-encoder U-Net with deep supervision, 40-patient SEGTHOR train set (corrected labels), median spacing + HU window + z-score + ROI crop, CE + soft Dice loss. 2.5D = the slice and its ±1 neighbours as 3 input channels, centre slice predicted



HD95 mm, 3D on the original CT grid, best checkpoint; each run's value is the mean over the patients of its validation fold. All organs = mean of the 4 organs per run. Y-axes differ per organ.